

Exercises - Class 3: Vectors

Sander Bossier

2023-10-21

1.

Create a 13x2 matrix `mat1` with the elements of the (built-in) vector `letters` in row-major order.

```
mat1 <- matrix(data = letters, ncol = 2, byrow = TRUE)
mat1
```

```
##      [,1] [,2]
## [1,] "a"  "b"
## [2,] "c"  "d"
## [3,] "e"  "f"
## [4,] "g"  "h"
## [5,] "i"  "j"
## [6,] "k"  "l"
## [7,] "m"  "n"
## [8,] "o"  "p"
## [9,] "q"  "r"
## [10,] "s" "t"
## [11,] "u" "v"
## [12,] "w" "x"
## [13,] "y" "z"
```

2.

Create another 13x2 matrix `mat2` with the elements of letters in column-major order. How can you transform `mat2` in order to make it identical to `mat1` (from the previous exercise)?

```
mat2 <- matrix(data = letters, nrow = 13)
mat2
```

```
##      [,1] [,2]
## [1,] "a"  "n"
## [2,] "b"  "o"
## [3,] "c"  "p"
## [4,] "d"  "q"
## [5,] "e"  "r"
## [6,] "f"  "s"
## [7,] "g"  "t"
## [8,] "h"  "u"
## [9,] "i"  "v"
## [10,] "j" "w"
## [11,] "k" "x"
## [12,] "l" "y"
## [13,] "m" "z"
```

```
identical(mat1, t(matrix(mat2, nrow = 2, ncol = 13)))
```

```
## [1] TRUE
```

3.

Extract the 2nd column from mat1 as a 13x1 matrix, not a vector. Do this for three different ways of indexing:
* positive integers * negative integers * logical vector

```
mat1[,2, drop = FALSE]
```

```
##      [,1]  
## [1,] "b"  
## [2,] "d"  
## [3,] "f"  
## [4,] "h"  
## [5,] "j"  
## [6,] "l"  
## [7,] "n"  
## [8,] "p"  
## [9,] "r"  
## [10,] "t"  
## [11,] "v"  
## [12,] "x"  
## [13,] "z"
```

```
mat1[,-1, drop = FALSE]
```

```
##      [,1]  
## [1,] "b"  
## [2,] "d"  
## [3,] "f"  
## [4,] "h"  
## [5,] "j"  
## [6,] "l"  
## [7,] "n"  
## [8,] "p"  
## [9,] "r"  
## [10,] "t"  
## [11,] "v"  
## [12,] "x"  
## [13,] "z"
```

```
mat1[,c(FALSE, TRUE), drop = FALSE]
```

```
##      [,1]  
## [1,] "b"  
## [2,] "d"  
## [3,] "f"  
## [4,] "h"  
## [5,] "j"  
## [6,] "l"  
## [7,] "n"  
## [8,] "p"  
## [9,] "r"  
## [10,] "t"
```

```
## [11,] "v"
## [12,] "x"
## [13,] "z"
```

4.

Compare the computation time of the three indexing operations from the previous exercise. You can use R's core function for timing but you will find the microbenchmark package more informative.

```
library(microbenchmark)
exp1 <- expression(mat1[,2, drop = FALSE])
exp2 <- expression(mat1[, -1, drop = FALSE])
exp3 <- expression(mat1[, c(FALSE, TRUE), drop = FALSE])
microbenchmark(exp1, exp2, exp3)
```

```
## Warning in microbenchmark(exp1, exp2, exp3): Could not measure a positive
## execution time for 20 evaluations.
```

```
## Unit: nanoseconds
##   expr min lq mean median uq   max neval
## exp1   0  0    2      0  0  200    100
## exp2   0  0   22      0  0 2000    100
## exp3   0  0    5      0  0  500    100
```

5.

Extract the five vowels from mat1 using:

- a single index vector
- an index matrix

Again, time the computations of both operations.

```
microbenchmark(
  mat1[c(1, 3, 5, 8, 11)],
  mat1[rbind(c(1, 1), c(3, 1), c(5, 1), c(8, 1), c(11, 1))], unit = "ns"
)
```

```
## Unit: nanoseconds
##                                     expr min   lq mean
##                                     mat1[c(1, 3, 5, 8, 11)] 400  500  837
## mat1[rbind(c(1, 1), c(3, 1), c(5, 1), c(8, 1), c(11, 1))] 3200 3500 4071
## median    uq   max neval
##      600  800 17500   100
##     3700 4100 31300   100
```

6.

What is the most efficient way to raise the first element of the following vector to the power 1, the second element to the power 2 etc. and the last element to the power 10 (you can copy-paste the following code):

```
vec1 <- 10:1
vec1
```

```
## [1] 10  9  8  7  6  5  4  3  2  1
```

```
vec1^(1:10)
```

```
## [1]    10    81   512  2401  7776 15625 16384  6561   512     1
```

7.

What is the most efficient way to raise the first element of the vector `vec1` (from the previous exercise) to the power 3, the second element to the power 2, the third element again to the power 3, the fourth element again to the power 2 etc. up until the one-but-last element to the power 3 and the last element to the power 2? In other words, your result should be: `[1] 1000 81 512 49 216 25 64 9 8 1`

```
vec1^rep(c(3,2), times = 5)
```

```
## [1] 1000 81 512 49 216 25 64 9 8 1
```

```
vec1^c(3,2)
```

```
## [1] 1000 81 512 49 216 25 64 9 8 1
```

The first elements of the shorter vector are consecutively copied to the end of the vector until its length matches that of the longer vector. **But** R raises a warning about this only when the length of the shorter vector is not a multiple of the length of the longer vector.

8.

Consider the following matrix (you can copy-paste the code):

```
mat3 <- matrix(seq(30, 270, by = 30), nrow = 3, ncol = 3)
mat3
```

```
##      [,1] [,2] [,3]
## [1,]   30  120  210
## [2,]   60  150  240
## [3,]   90  180  270
```

```
mat3/c(2,3,5)
```

```
##      [,1] [,2] [,3]
## [1,]   15   60  105
## [2,]   20   50   80
## [3,]   18   36   54
```

```
mat3/rep(c(2,3,5), each = 3)
```

```
##      [,1] [,2] [,3]
## [1,]   15  40  42
## [2,]   30  50  48
## [3,]   45  60  54
```

```
t(t(mat3)/c(2,3,5))
```

```
##      [,1] [,2] [,3]
## [1,]   15  40  42
## [2,]   30  50  48
## [3,]   45  60  54
```

```
sweep(x = mat3, MARGIN = 1, STATS = c(2, 3, 5), FUN = "/")
```

```
##      [,1] [,2] [,3]
## [1,]   15  60  105
## [2,]   20  50   80
## [3,]   18  36   54
```

```
sweep(x = mat3, MARGIN = 2, STATS = c(2, 3, 5), FUN = "/")
```

```
##      [,1] [,2] [,3]
## [1,]   15  40  42
## [2,]   30  50  48
## [3,]   45  60  54
```

9.

Use a meta-function to generate all possible sums of the sides of two dice. The sides of one die can be represented as a vector (you can copy-paste the code):

```
vec2 <- 1:6
vec2
```

```
## [1] 1 2 3 4 5 6
```

```
outer(vec2, vec2, FUN = "+")
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]    2    3    4    5    6    7
## [2,]    3    4    5    6    7    8
## [3,]    4    5    6    7    8    9
## [4,]    5    6    7    8    9   10
## [5,]    6    7    8    9   10   11
## [6,]    7    8    9   10   11   12
```

10.

The matrix from the previous exercise can be converted to a vector. There is another meta-function which allows you to compute the vector of all (36) sums of the sides of two dice. Compare the computation times of both commands.

```
vec3 <- as.vector(outer(vec2, vec2, FUN = "+"))
vec3
```

```
## [1] 2 3 4 5 6 7 3 4 5 6 7 8 4 5 6 7 8 9 5 6 7 8 9 10 6
## [26] 7 8 9 10 11 7 8 9 10 11 12
```

```
microbenchmark(as.vector(outer(vec2, vec2, FUN = "+")), unit = "ns")
```

```
## Unit: nanoseconds
```

```
##              expr   min    lq  mean median    uq   max
## as.vector(outer(vec2, vec2, FUN = "+")) 8200 8500 10314   8800 9200 113600
## neval
##      100
```

```
microbenchmark(apply(expand.grid(vec2, vec2), MARGIN = 1, FUN = sum), unit = "ns")
```

```
## Unit: nanoseconds
```

```
##              expr   min    lq  mean
## apply(expand.grid(vec2, vec2), MARGIN = 1, FUN = sum) 173700 180950 239080
## median      uq      max neval
## 186500 200650 1393600   100
```

11.

Compare the computation time of R's core `cumsum()` function with `Reduce(. , accumulate = TRUE)`. Do this on this `vec3` (you can copy-paste the code):

```

set.seed(123)
vec3 <- sample(vec1, size = 100, replace = TRUE)
head(vec3, n = 10)

## [1] 8 8 1 9 5 6 7 5 2 1

tail(vec3, n = 10)

## [1] 5 8 3 8 3 10 4 4 4 1

microbenchmark(cumsum(vec3), times = 100, unit = "ns")

## Unit: nanoseconds
##      expr min  lq mean median  uq   max neval
## cumsum(vec3) 200 300  741    300 500 23900   100

microbenchmark(Reduce(f = "+", x = vec3, accumulate = TRUE), times = 100, unit = "ns")

## Unit: nanoseconds
##      expr      min      lq mean median      uq
## Reduce(f = "+", x = vec3, accumulate = TRUE) 59500 60850 66431  62500 65250
##      max neval
## 178600   100

```

12.

Compare the computation times of three ways to sum every two consecutive elements in vector `vec3` (from the previous exercise):

- a for loop
- the `vapply()` meta-function
- another meta-function of the `apply` family, closely related to `vapply()`

In other words, the first element of your result 16 should be the sum of the first two elements 8 and 8 in `vec3`, the second element of the result 9 should be the sum of the second element 8 and third element 1, the third element of the result 10 needs to be the sum of the third element 1 and fourth element 9 in `vec3`, and so on:

```

[1] 16 9 10 14 11 13 12 7 3 7 14 10 4 4 10 11 4 5 5 3 10 15 17 14 10
[26] 7 5 6 4 3 5 10 10 10 11 7 11 15 9 12 19 12 4 7 11 8 3 8 12 8
[51] 8 10 9 14 15 14 19 19 16 13 11 13 10 9 12 7 4 6 12 11 5 10 12 12 12
[76] 9 6 7 12 9 11 9 10 10 10 10 6 12 17 15 13 11 11 11 13 14 8 8 5

```

This exercise assumes some basic familiarity with creating functions in R.

```

consecsumloop <- function(vec) {
  result <- numeric(length(vec) - 1)
  for (i in 1:(length(vec) - 1)) {
    result[i] <- vec[i] + vec[i+1]
  }
  result
}

microbenchmark({result <- numeric(length(vec3) - 1)
  for (i in 1:(length(vec3) - 1)) {
    result[i] <- vec3[i] + vec3[i+1]
  }

```

```

}
result}, unit = "us")

## Unit: microseconds
##
## {      result <- numeric(length(vec3) - 1)      for (i in 1:(length(vec3) - 1)) {      result[i] <-
##      min      lq      mean median      uq      max neval
## 2985.9 3195.55 4163.804 3344.25 3867.3 20144.2    100
microbenchmark(
  vapply(seq_len(length(vec3) - 1), FUN = function(i) sum(vec3[i:i + 1]), numeric(1)), unit = "us"
)

## Unit: microseconds
##
##      vapply(seq_len(length(vec3) - 1), FUN = function(i) sum(vec3[i:i +      1]), numeric(1))      expr
##      min      lq      mean median      uq      max neval
## 104.1 107.7 154.57 111.45 120.85 2508    100
microbenchmark(
  sapply(seq_len(length(vec3) - 1), function(i) sum(vec3[i:i + 1])), unit = "us"
)

## Unit: microseconds
##
##      sapply(seq_len(length(vec3) - 1), function(i) sum(vec3[i:i +      1]))      expr      min
##      lq      mean median      uq      max neval
## 121 161.271 123.05 129.8 2348    100

```