

# Exercises - Class 4: Functions

Sander Bossier

2023-10-22

## 1.

Harmonic numbers are sums of the reciprocals of the positive integers. More specifically, the  $n$ 'th harmonic number is:

$$\sum_{i=1}^n \frac{1}{i}$$

Create a function `Harmonic(n)` which computes the  $n$ th harmonic number. The body of this function can consist of a single line. An example is:

```
Harmonic <- function(n) sum(1 / (1:n))
```

```
Harmonic(5)
```

```
## [1] 2.283333
```

## 2.

Harmonic numbers are recursive:

$$H_n = H_{n-1} + \frac{1}{n}$$

Write a recursive function `rHarmonic(n)` which computes the  $n$ th harmonic number recursively (in other words, the result should be the same as in the previous exercise).

```
rHarmonic <- function(n) {  
  if (n == 1) {  
    1  
  }  
  else{  
    rHarmonic(n - 1) + (1 / n)  
  }  
}  
rHarmonic(5)
```

```
## [1] 2.283333
```

## 3.

Compare the computation time of `Harmonic(n)` and `rHarmonic(n)` for the 1000th harmonic number (using R's core timing function).

```
system.time(  
  Harmonic(1000)  
)
```

```
##      user  system elapsed
##         0         0         0
```

```
system.time(
  rHarmonic(1000)
)
```

```
##      user  system elapsed
##    0.00    0.00    0.01
```

```
system.time(
  Harmonic(1e6)
)
```

```
##      user  system elapsed
##    0.00    0.00    0.02
```

#### 4.

Create a binary infix operator `%fill%` which replaces all `NA` values in the vector on the left-hand side by the value specified on the right-hand side.

```
`%fill%` <- function(vec, value) {
  vec[is.na(vec)] <- value # or
  #replace(vec, is.na(vec), value)
  vec
}

vec1 <- c(5, 55, NA, -9, -99, NA, NA, -5, -9, 999)
vec2 <- vec1 %fill% 0
vec2
```

```
## [1]  5 55  0 -9 -99  0  0 -5 -9 999
```

#### 5.

Create a replacement function `fill()` which replaces all `NA` values in a vector by the value specified on the right-hand side.

```
`fill` <- function(vec, value) {
  vec[is.na(vec)] <- value
  vec
}

vec1 <- c(5, 55, NA, -9, -99, NA, NA, -5, -9, 999)
fill(vec1) <- 0
vec1
```

```
## [1]  5 55  0 -9 -99  0  0 -5 -9 999
```

#### 6.

Compare the computation time of the `%fill%` operator and the `fill()` replacement function on the following large vector (you can copy-paste these two lines):

```
library(microbenchmark)
set.seed(123)
vec3 <- sample(c(5, -9, NA), size = 1e6, replace = TRUE)
```

Do this using both R's core function and the package `microbenchmark` (replace the NA's by 0).

```
system.time(  
  tmp <- vec3 %fill% 0  
)
```

```
##    user  system elapsed  
##    0.00    0.00    0.03
```

```
system.time(  
  fill(vec3) <- 0  
)
```

```
##    user  system elapsed  
##    0.00    0.00    0.02
```

```
microbenchmark(  
  tmp <- vec3 %fill% 0,  
  unit = "us"  
)
```

```
## Unit: microseconds  
##           expr      min       lq      mean  median       uq      max neval  
## tmp <- vec3 %fill% 0 3545.801 3722.901 4869.805 4041.9 5133.952 9220.201   100
```

```
microbenchmark(  
  fill(vec3) <- 0,  
  unit = "us"  
)
```

```
## Unit: microseconds  
##           expr      min       lq      mean  median       uq      max neval  
## fill(vec3) <- 0 3551.102 3719.851 4986.207 4208.801 5252.801 10261.7   100
```

## 7.

Create two functions `full1()` and `full2()` which both extract the non-NA elements from a vector. One of both functions needs to make use of the functional `Filter()` (see last week). Compare their computation times on the `vec3` object from the previous exercises. You can use both R's core function and the `microbenchmark` package but for the latter you can specify a number of `times <= 100`.

Note: There is even a way to combine the `Filter()` function with the `Negate()` function. If you manage to find that one, then you can compare all three functions.