

Exercises - Class 4: Functions

Sander Bossier

2023-10-22

1.

Harmonic numbers are sums of the reciprocals of the positive integers. More specifically, the n 'th harmonic number is:

$$\sum_{i=1}^n \frac{1}{i}$$

Create a function `Harmonic(n)` which computes the n th harmonic number. The body of this function can consist of a single line. An example is:

```
Harmonic <- function(n) sum(1 / (1:n))
```

```
Harmonic(5)
```

```
## [1] 2.283333
```

2.

Harmonic numbers are recursive:

$$H_n = H_{n-1} + \frac{1}{n}$$

Write a recursive function `rHarmonic(n)` which computes the n th harmonic number recursively (in other words, the result should be the same as in the previous exercise).

```
rHarmonic <- function(n) {  
  if (n == 1) {  
    1  
  }  
  else{  
    rHarmonic(n - 1) + (1 / n)  
  }  
}  
rHarmonic(5)
```

```
## [1] 2.283333
```

3.

Compare the computation time of `Harmonic(n)` and `rHarmonic(n)` for the 1000th harmonic number (using R's core timing function).

```
system.time(  
  Harmonic(1000)  
)
```

```
##      user  system elapsed
##         0         0         0
```

```
system.time(
  rHarmonic(1000)
)
```

```
##      user  system elapsed
##         0         0         0
```

```
system.time(
  Harmonic(1e6)
)
```

```
##      user  system elapsed
##    0.03    0.00    0.00
```

4.

Create a binary infix operator `%fill%` which replaces all `NA` values in the vector on the left-hand side by the value specified on the right-hand side.

```
`%fill%` <- function(vec, value) {
  vec[is.na(vec)] <- value # or
  #replace(vec, is.na(vec), value)
  vec
}

vec1 <- c(5, 55, NA, -9, -99, NA, NA, -5, -9, 999)
vec2 <- vec1 %fill% 0
vec2
```

```
## [1]  5  55  0 -9 -99  0  0 -5 -9 999
```

5.

Create a replacement function `fill()` which replaces all `NA` values in a vector by the value specified on the right-hand side.

```
`fill` <- function(vec, value) {
  vec[is.na(vec)] <- value
  vec
}

vec1 <- c(5, 55, NA, -9, -99, NA, NA, -5, -9, 999)
fill(vec1) <- 0
vec1
```

```
## [1]  5  55  0 -9 -99  0  0 -5 -9 999
```

6.

Compare the computation time of the `%fill%` operator and the `fill()` replacement function on the following large vector (you can copy-paste these two lines):

```
library(microbenchmark)
set.seed(123)
vec3 <- sample(c(5, -9, NA), size = 1e6, replace = TRUE)
```

Do this using both R's core function and the package `microbenchmark` (replace the NA's by 0).

```
system.time(  
  tmp <- vec3 %fill% 0  
)
```

```
##    user  system elapsed  
##    0.00    0.00    0.02
```

```
system.time(  
  fill(vec3) <- 0  
)
```

```
##    user  system elapsed  
##    0.01    0.00    0.03
```

```
microbenchmark(  
  tmp <- vec3 %fill% 0,  
  unit = "us"  
)
```

```
## Unit: microseconds  
##           expr   min      lq     mean   median      uq     max neval  
## tmp <- vec3 %fill% 0 3543 3694.65 4880.555 4207.551 5128.851 11381.2   100
```

```
microbenchmark(  
  fill(vec3) <- 0,  
  unit = "us"  
)
```

```
## Unit: microseconds  
##           expr   min      lq     mean   median      uq     max neval  
## fill(vec3) <- 0 3542.601 3744.651 4958.441 4363.201 5046.901 12335.3   100
```

7.

Create two functions `full1()` and `full2()` which both extract the non-NA elements from a vector. One of both functions needs to make use of the functional `Filter()` (see last week). Compare their computation times on the `vec3` object from the previous exercises. You can use both R's core function and the **`microbenchmark`** package but for the latter you can specify a number of `times <= 100`.

Note: There is even a way to combine the `Filter()` function with the `Negate()` function. If you manage to find that one, then you can compare all three functions.

```
full1 <- function(vec) {  
  vec[!is.na(vec)]  
}  
  
full2 <- function(vec) {  
  Filter(function(x) !is.na(x), vec)  
}  
  
full3 <- function(vec) {  
  Filter(Negate(function(x) is.na(x)), vec)  
}
```

```
vec1 <- c(5, 55, NA, -9, -99, NA, NA, -5, -9, 999)  
microbenchmark(  
  full1(vec3),
```

```

full12(vec3),
full13(vec3),
times = 100,
unit = "ns"
)

## Unit: nanoseconds
##      expr      min      lq      mean      median      uq
## full1(vec3) 9472500 14876601 23071048 17656051 24042951
## full2(vec3) 553414001 913735801 1109730769 1024485602 1230508601
## full3(vec3) 967064801 1307615051 221269689718 1423236452 1645245701
##      max neval
## 1.044511e+08 100
## 1.856980e+09 100
## 2.160360e+13 100

```

8.

One of the features of R's syntax is that a function name before the opening bracket (can actually be a variable. This allows uses such as the following:

```

f <- sum
f(vec2)

```

```
## [1] 937
```

```

f <- mean
f(vec2)

```

```
## [1] 93.7
```

Create a function `twofun()` which has two formal arguments:

- `x` which can be a numeric vector
- `fun` which can be any function computing some descriptive statistic

The function `twofun()` should apply the function specified as `fun` to every two consecutive elements in the vector specified as `x`. In other words, if `fun = sum` then `twofun()` will compute the sum of every two consecutive elements in `x` like we have seen last week (see the last exercise of the previous week). However, `twofun()` allows for more flexibility:

```

twofun <- function(vec, fun = sum) {
  vapply(
    seq_len(length(vec) - 1),
    FUN = function(i) fun(vec[i:(i + 1)]),
    numeric(1)
  )
}

```

```
twofun(vec1, fun = sum)
```

```
## [1] 60 NA NA -108 NA NA NA -14 990
```

```
twofun(vec1, fun = mean)
```

```
## [1] 30 NA NA -54 NA NA NA -7 495
```

```
twofun(vec1, fun = median)
```

```
## [1] 30 NA NA -54 NA NA NA -7 495
```

```
twofun(vec1, fun = var)

## [1] 1250 NA NA 4050 NA NA NA 8 508032
twofun(vec1)

## [1] 60 NA NA -108 NA NA NA -14 990
```

9.

R's functions for descriptive statistics have an argument `na.rm` and `var()` even has an argument `use`. Rewrite the `twosum()` function so that it can handle all of these extra arguments without any conflict.

```
twofun <- function(vec, fun = sum, ...) {

  vapply(
    X = seq_len(length(vec) - 1),
    FUN = function(i) fun(vec[i:(i + 1)], ...),
    FUN.VALUE = numeric(1)
  )
}
```

10.

```
twofunfun <- function(fun = sum) {
  out <- function(vec, ...) {
    vapply(
      X = seq_len(length(vec) - 1),
      FUN = function(i) fun(vec[i:(i + 1)], ...),
      FUN.VALUE = numeric(1)
    )
  }
  out
}
```

```
twomean <- twofunfun(fun = mean)
twomean(vec1)

## [1] 30 NA NA -54 NA NA NA -7 495
twovar <- twofunfun(fun = var)
twovar(vec1)

## [1] 1250 NA NA 4050 NA NA NA 8 508032
environment(twomean)

## <environment: 0x00000258c5b78538>
environment(twovar)

## <environment: 0x00000258c59f6e28>
exists("fun", environment(twomean))

## [1] TRUE
get("fun", environment(twomean))
```

```

## function (x, ...)
## UseMethod("mean")
## <bytecode: 0x00000258c48b1598>
## <environment: namespace:base>
exists("fun", environment(twoovar))

## [1] TRUE
get("fun", environment(twoovar))

## function (x, y = NULL, na.rm = FALSE, use)
## {
##   if (missing(use))
##     use <- if (na.rm)
##       "na.or.complete"
##     else "everything"
##   na.method <- pmatch(use, c("all.obs", "complete.obs", "pairwise.complete.obs",
##     "everything", "na.or.complete"))
##   if (is.na(na.method))
##     stop("invalid 'use' argument")
##   if (is.data.frame(x))
##     x <- as.matrix(x)
##   else stopifnot(is.atomic(x))
##   if (is.data.frame(y))
##     y <- as.matrix(y)
##   else stopifnot(is.atomic(y))
##   .Call(C_cov, x, y, na.method, FALSE)
## }
## <bytecode: 0x00000258c260f6b8>
## <environment: namespace:stats>

```